

Oct 07, 2025 23:09

**pyproject.toml**

Page 1/1

```
# pyproject.toml.jinja

[project]
name = "hw04"
version = "0.1.0"
description = "CIFAR10"
readme = "README.md"
authors = [
    { name = "Josh Miao", email = "josh.miao@cooper.edu" }
]
requires-python = ">=3.13"
dependencies = [
    "structlog",
    "rich",
    "numpy",
    "jax[cuda12]",
    "flax",
    "typing>=3.10.0.0",
    "tqdm>=4.67.1",
    "pydantic>=2.11.9",
    "pydantic-settings>=2.11.0",
    "optax>=0.2.6",
    "dataclasses>=0.8",
    "tensorflow>=2.20.0",
    "nnx>=0.0.8",
    "diffusers==0.11.1",
    "orbax>=0.1.9",
    "orbax-checkpoint>=0.11.25",
    "pathlib>=1.0.1",
    "path>=17.1.1",
]

[project.scripts]
hw04 = "hw04:main"
test10 = "hw04:test10"
test100 = "hw04:test100"

[build-system]
requires = ["uv_build>=0.8.3,<0.9.0"]
build-backend = "uv_build"
```

Oct 09, 2025 2:39                    config.py                    Page 1/2

```

from pathlib import Path
from importlib.resources import files
from typing import Tuple, List

from pydantic import BaseModel
from pydantic_settings import (
    BaseSettings,
    PydanticBaseSettingsSource,
    SettingsConfigDict,
    TomlConfigSettingsSource,
)

class DataSettings(BaseModel):
    """Settings for data generation."""
    train_ratio: float = 0.8

    #in data.py, change back to cifar10 and change num_classes here in model seeting
    s back to 10
class ModelSettings(BaseModel):
    input_depth: int = 3
    num_classes: int = 10
    layer_depth: List[int] = [64, 128, 256]
    layer_kernel_sizes: List[List[int]] = [[3, 3], [3,3], [3,3]]
    strides: List[int] = [1, 1, 1]
    #dropout_rate: float = 0.2
    shape: List[int] = [32, 32]
    noise_std: int = 0.05

class TrainingSettings(BaseModel):
    """Settings for model training."""
    batch_size: int = 256
    num_iters: int = 150000
    learning_rate: float = 1.5e-4
    L2_weight: float = 1.5e-4
    #tried to implement cosine decay with linear warmup but that was low key che
eks
    #warmup_steps: int = 800
    #initial_rate: float = 1e-5
    #end_value: float = 1e-5

class AppSettings(BaseSettings):
    """Main application settings."""

    debug: bool = False
    random_seed: int = 427 #for ece 427 ofc
    data: DataSettings = DataSettings()
    training: TrainingSettings = TrainingSettings()
    model: ModelSettings = ModelSettings()

    model_config = SettingsConfigDict(
        toml_file=files("hw04").joinpath("config.toml"),
        env_nested_delimiter="_",
    )

    @classmethod
    def settings_customise_sources(
        cls,
        settings_cls: type[BaseSettings],
        init_settings: PydanticBaseSettingsSource,
        env_settings: PydanticBaseSettingsSource,
        dotenv_settings: PydanticBaseSettingsSource,
        file_secret_settings: PydanticBaseSettingsSource,
    ) -> tuple[PydanticBaseSettingsSource, ...]:
        """
        Set the priority of settings sources.
        We use a TOML file for configuration.
        """
        return (

```

Oct 09, 2025 2:39                    config.py                    Page 2/2

```

        init_settings,
        TomlConfigSettingsSource(settings_cls),
        env_settings,
        dotenv_settings,
        file_secret_settings,
    )

def load_settings() -> AppSettings:
    """Load application settings."""
    return AppSettings()

```

Oct 09, 2025 9:34

## data.py

Page 1/1

```

from dataclasses import dataclass, InitVar, field
import numpy as np
import structlog
import tensorflow as tf

log = structlog.get_logger()

@dataclass
class cifar10:
    """retrive data. also i know that my class name is misleading when im training cifar10
    kinda too lazy to change it."""
    rng: InitVar[np.random.Generator]
    train_ratio: float = 0.8

    # Data arrays (initialized post-construction)
    x_train: np.ndarray = field(init=False)
    x_val: np.ndarray = field(init=False)
    x_test: np.ndarray = field(init=False)
    y_train: np.ndarray = field(init=False)
    y_val: np.ndarray = field(init=False)
    y_test: np.ndarray = field(init=False)

    index: np.ndarray = field(init=False)

    def __post_init__(self, rng:np.random.Generator):
        # Load and normalize cifar10 data
        (x_train_raw, y_train_raw), (x_test_raw, y_test_raw) = tf.keras.datasets.cifar10.load_data()

        # convert to float32 in [0,1]
        x_train_raw = x_train_raw.astype(np.float32) / 255.0
        x_test_raw = x_test_raw.astype(np.float32) / 255.0

        # labels come as shape (N,1) flatten to (N,)
        y_train_raw = y_train_raw.reshape(-1)
        y_test_raw = y_test_raw.reshape(-1)

        # shuffle and split training into train/val
        train_idx, val_idx = self.split_indices(len(y_train_raw), rng)

        self.x_train = x_train_raw[train_idx]
        self.y_train = y_train_raw[train_idx]
        self.x_val = x_train_raw[val_idx]
        self.y_val = y_train_raw[val_idx]
        self.x_test = x_test_raw
        self.y_test = y_test_raw
        self.index = np.arange(len(self.y_train))

    def split_indices(self, size: int, rng: np.random.Generator):
        indices = np.arange(size)
        rng.shuffle(indices)
        split = int(self.train_ratio * size)
        return indices[:split], indices[split:]

    def get_batch(self, rng: np.random.Generator, batch_size: int):
        choices = rng.choice(self.index, size=batch_size)
        return self.x_train[choices], self.y_train[choices]

    def get_validation(self):
        return self.x_val, self.y_val

    def get_test(self):
        return self.x_test, self.y_test

```

Sep 29, 2025 22:46

**logging.py**

Page 1/1

```
import logging
import os
import sys
from pathlib import Path

import jax
import numpy as np
import structlog

class FormattedFloat(float):
    def __repr__(self) -> str:
        return f"{self:.4g}"


def custom_serializer_processor(logger, method_name, event_dict):
    for key, value in event_dict.items():
        # Handle JAX arrays in addition to TF tensors
        if hasattr(value, "numpy"): # Covers TF tensors
            value = value.numpy()
        if isinstance(value, jax.Array):
            value = np.array(value)
        if isinstance(value, (np.generic, np.ndarray)):
            value = value.item() if value.size == 1 else value.tolist()
        if isinstance(value, float):
            value = FormattedFloat(value)
        if isinstance(value, Path):
            value = str(value)
        event_dict[key] = value
    return event_dict


def configure_logging():
    """Configure logging for the application."""
    logging.basicConfig(
        format="%(message)s",
        stream=sys.stdout,
    )

    # Set the level for the application's logger
    log_level = os.environ.get("LOG_LEVEL", "INFO").upper()
    logging.getLogger("hw04").setLevel(log_level)

    structlog.configure(
        processors=[
            structlog.stdlib.filter_by_level,
            structlog.stdlib.add_logger_name,
            structlog.stdlib.add_log_level,
            structlog.stdlib.PositionalArgumentsFormatter(),
            structlog.processors.TimeStamper(fmt="iso"),
            structlog.processors.StackInfoRenderer(),
            structlog.processors.format_exc_info,
            structlog.processors.UnicodeDecoder(),
            custom_serializer_processor,
            structlog.dev.ConsoleRenderer(
                colors=True, exception_formatter=structlog.dev.RichTracebackFormatter
            ),
        ],
        context_class=dict,
        logger_factory=structlog.stdlib.LoggerFactory(),
        wrapper_class=structlog.stdlib.BoundLogger,
        cache_logger_on_first_use=True,
    )
```

Oct 05, 2025 18:43

model.py

Page 1/5

```

import jax
import jax.numpy as jnp
from flax import nnx
from typing import Union, Tuple, List
import structlog
import math

from .config import ModelSettings

settings = ModelSettings()

log = structlog.get_logger()

#Pooling
class Pooling(nn.Module):
    def __init__(self, *, window=(2,2), strides=(2,2), padding="SAME"):
        self.window = window
        self.strides = strides
        self.padding = padding

    def __call__(self, x, train=False):
        return jax.lax.reduce_window(
            x,
            -jnp.inf,
            jax.lax.max,
            window_dimensions=(1, self.window[0], self.window[1], 1),
            window_strides=(1, self.strides[0], self.strides[1], 1),
            padding=self.padding,
        )

    # GroupNorm
class GroupNorm(nn.Module):
    """
    Group normalization with learnable affine parameters per channel.
    """
    def __init__(self, *, num_channels: int, num_groups: int = 8, epsilon: float = 1e-5):
        # store simple config
        self.num_channels = int(num_channels)
        self.num_groups = int(num_groups)
        self.epsilon = float(epsilon)

        # choose group count that divides channels (fallback by decrementing)
        G = min(self.num_groups, self.num_channels)
        while G > 1 and (self.num_channels % G) != 0:
            G -= 1
        self._G = max(1, G)

        # learnable affine parameters (per-channel)
        # nnx.Param wraps a leaf parameter; we initialize to gamma=1, beta=0
        self.scale = nnx.Param(jnp.ones((self.num_channels,), dtype=jnp.float32))
        self.bias = nnx.Param(jnp.zeros((self.num_channels,), dtype=jnp.float32))

    def __call__(self, x: jnp.ndarray, train: bool = False) -> jnp.ndarray:
        # x: (N, H, W, C)
        N, H, W, C = x.shape
        assert C == self.num_channels, "GroupNorm: channel mismatch"

        # reshape into groups: (N, H, W, G, C_per_group)
        Cg = C // self._G
        xg = x.reshape((N, H, W, self._G, Cg))

        # compute mean/var over (H, W, Cperg)
        mean = jnp.mean(xg, axis=(1, 2, 4), keepdims=True)
        var = jnp.var(xg, axis=(1, 2, 4), keepdims=True)

        xg = (xg - mean) / jnp.sqrt(var + self.epsilon)

```

Oct 05, 2025 18:43

model.py

Page 2/5

```

x_norm = xg.reshape((N, H, W, C))

# apply affine per-channel (broadcast over N,H,W)
gamma = self.scale.value if hasattr(self.scale, "value") else self.scale
beta = self.bias.value if hasattr(self.bias, "value") else self.bias
gamma = gamma.reshape((1, 1, 1, C))
beta = beta.reshape((1, 1, 1, C))

return x_norm * gamma + beta

#Convolution
class Convolution(nn.Module):
    """
    Single convolution with L2
    """
    def __init__(self,
                 *,
                 keys,
                 L2_weight: float,
                 in_features: int,
                 out_features: int,
                 kernel_size: Union[int, Tuple[int, int]],
                 strides: Union[int, Tuple[int, int]],
                 padding: str = "SAME",
                 ):
        # store config
        self.L2_weight = float(L2_weight)
        self.padding = padding

        # normalize kernel and strides
        if isinstance(kernel_size, int):
            kh = kw = int(kernel_size)
        else:
            kh, kw = map(int, kernel_size)

        if isinstance(strides, int):
            sh = sw = int(strides)
        else:
            sh, sw = strides

        self.kernel_size = (kh, kw)
        self.strides = (sh, sw)

        # compute fan_in and He (Kaiming) std for ReLU
        fan_in = kh * kw * int(in_features)
        he_std = math.sqrt(2.0 / float(fan_in))

        # generate kernel
        key = keys if keys is not None else jax.random.PRNGKey(0)
        kernel_shape = (kh, kw, int(in_features), int(out_features))
        init_kernel = jax.random.normal(key, kernel_shape, dtype=jnp.float32) *
        jnp.array(he_std, dtype=jnp.float32)
        self.kernel = nnx.Param(init_kernel)

    def __call__(self, x: jnp.ndarray, train: bool = False) -> jnp.ndarray:
        """
        x expected shape: (N, H, W, C) (NHWC).
        Returns NHWC.
        """
        kernel_arr = self.kernel.value if hasattr(self.kernel, "value") else self.kernel

        out = jax.lax.conv_general_dilated(
            lhs=x,
            rhs=kernel_arr,
            window_strides=self.strides,
            padding=self.padding,
            dimension_numbers=("NHWC", "HWIO", "NHWC"),

```

Oct 05, 2025 18:43

## model.py

Page 3/5

```

        )
    return out

def L2_loss(self) -> jnp.ndarray:
    k = self.kernel.value if hasattr(self.kernel, "value") else self.kernel
    return jnp.sum(jnp.square(k)) * self.L2_weight


class GaussianNoise(nn.Module):
    def __init__(self, *, std: float, key):
        self.std = std
        self.rngs = nnx.Rngs(noise=key)

    def __call__(self, x, train: bool = False):
        if not train or self.std <= 0:
            return x
        noise = jax.random.normal(self.rngs.noise(), x.shape) * self.std
        return x + noise

# ResidualBlock
class ResidualBlock(nn.Module):
    """
    Basic residual block: Conv -> GroupNorm -> ReLU -> Conv -> GroupNorm -> skip -> ReLU
    """
    def __init__(self,
                *,
                keys, # tuple of three keys: (conv1_key, conv2_key, skip_proj_key)
                L2_weight: float,
                in_features: int,
                out_features: int,
                kernel_size: Tuple[int, int] = (3, 3),
                strides: Tuple[int, int] = (1, 1),
                ):
        keys, conv1_key, conv2_key, proj_key = keys

        # conv1: in_features -> out_features
        self.conv1 = Convolution(
            keys=conv1_key,
            L2_weight=L2_weight,
            in_features=in_features,
            out_features=out_features,
            kernel_size=kernel_size,
            strides=strides,
        )
        self.norm1 = GroupNorm(num_channels=out_features)

        # conv2: out_features -> out_features
        self.conv2 = Convolution(
            keys=conv2_key,
            L2_weight=L2_weight,
            in_features=out_features,
            out_features=out_features,
            kernel_size=kernel_size,
            strides=(1, 1),
        )
        self.norm2 = GroupNorm(num_channels=out_features)

        # projection on skip if needed
        if in_features != out_features or strides != (1, 1):
            self.proj = nnx.Conv(
                in_features=in_features,
                out_features=out_features,
                kernel_size=(1, 1),
                strides=strides,
                padding="SAME",
                rngs=nnx.Rngs(params=proj_key),
            )
        else:
    
```

Oct 05, 2025 18:43

## model.py

Page 4/5

```

        self.proj = None

        self.noise = GaussianNoise(std=settings.noise_std, key = proj_key)

    def __call__(self, x: jnp.ndarray, train: bool = False) -> jnp.ndarray:
        identity = x

        out = self.conv1(x, train)
        out = self.norm1(out, train)
        out = jax.nn.relu(out)

        if train:
            out = self.noise(out,train)

        out = self.conv2(out, train)
        out = self.norm2(out, train)

        if self.proj is not None:
            identity = self.proj(identity)

        out = out + identity
        return jax.nn.relu(out)

def L2_loss(self) -> jnp.ndarray:
    loss = self.conv1.L2_loss() + self.conv2.L2_loss()
    if self.proj is not None:
        loss += jnp.sum(jnp.square(self.proj.kernel.value))
    return loss

# classify
class Classify(nn.Module):
    """
    CNN classifier built from ResidualBlocks.
    """
    def __init__(self,
                *,
                input_depth: int,
                layer_depths: List[int],
                layer_kernel_sizes: List[Tuple[int, int]],
                strides: List[Union[int, Tuple[int, int]]],
                num_classes: int,
                L2_weight: float,
                rngs: nnx.Rngs,
                shape: List[int],
                ):
        keys = rngs.params() # root key (jax.Array)

        input_features = [input_depth] + layer_depths[:-1]
        layers = []

        # We'll need 3 keys per residual block plus one key for final linear:
        total_keys = len(layer_depths) * 3 + 1
        all_keys = jax.random.split(keys, total_keys)
        block_keys = all_keys[:-1]
        final_key = all_keys[-1]

        for i, out_ch in enumerate(layer_depths):
            k1, k2, k3 = block_keys[3*i:3*i+3]
            stride_val = (strides[i], strides[i]) if isinstance(strides[i], int)
        else strides[i]:
            block = ResidualBlock(
                keys=(k1, k2, k3),
                L2_weight=L2_weight,
                in_features=input_features[i],
                out_features=out_ch,
                kernel_size=tuple(layer_kernel_sizes[i]),
                strides=stride_val,
            )
    
```

Oct 05, 2025 18:43

## model.py

Page 5/5

```
layers.append(block)

# pooling between blocks (except after last)
if i < len(layer_depths) - 1:
    layers.append(Pooling(window=(2,2), strides=(2,2)))

self.layers = nnx.List(layers)

# compute flatten size after downsampling
h, w = shape[0], shape[1]
for s in strides:
    sh, sw = (s, s) if isinstance(s, int) else s
    h = math.ceil(h / sh)
    w = math.ceil(w / sw)
# account for pooling between blocks
h = h // (2 ** (len(layer_depths) - 1))
w = w // (2 ** (len(layer_depths) - 1))
flatten_size = h * w * layer_depths[-1]

# final linear
self.final_layer = nnx.Linear(
    in_features=flatten_size,
    out_features=num_classes,
    rngs=nnx.Rngs(params=final_key),
)

def __call__(self, x: jnp.ndarray, train: bool = False) -> jnp.ndarray:
    out = x
    for layer in self.layers:
        # pooling and blocks all accept (x, train) signature
        out = layer(out, train)
    out = out.reshape((out.shape[0], -1))
    return self.final_layer(out)

def L2_loss(self) -> jnp.ndarray:
    return sum(layer.L2_loss() for layer in self.layers if hasattr(layer, "L2_loss"))
```

Oct 09, 2025 9:03

## output-cifar10-final-final.txt

Page 1/1

```

==> Running homework 'hw04'...
/mnt/c/Users/joshu/OneDrive/Cooper Union/Deep Learning/hw04/.venv/lib/python3.13
/site-packages/pydantic/main.py:463: UserWarning: Pydantic serializer warnings:
  PydanticSerializationUnexpectedValue(Expected 'int' - serialized value may not
  be as expected [input_value=0.05, input_type=float])
  return self._pydantic_serializer_.to_python()
INFO:hw04:[2m2025-10-09T06:40:23.690409Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
mSettings loaded          ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m ^[[36msett
ings'^[[0m='[[35m{'debug': False, 'random_seed': 427, 'data': {'train_ratio': 0.8
}, 'training': {'batch_size': 256, 'num_iters': 150000, 'learning_rate': 0.00015
, 'L2_weight': 0.00015}, 'model': {'input_depth': 3, 'num_classes': 10, 'layer_d
epth': [64, 128, 256], 'layer_kernel_sizes': [[3, 3], [3, 3], [3, 3]], 'strides'
: [1, 1, 1], 'shape': [32, 32], 'noise_std': 0.05}}^[[0m
INFO:hw04:[2m2025-10-09T06:40:28.617145Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
mcifar10 data retrieved      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m ^[[36mshap
e^[[0m='[[35m(40000, 32, 32, 3)^[[0m
INFO:hw04.training:^[[2m2025-10-09T06:40:52.363057Z^[[0m [[[[32m^[[1m[info      ^[
[0m] ^[[1mstarted training      ^[[0m [[[[0m^[[1m^[[34mhw04.training^[[0
m]^[[0m

 0% | 0/150000 [00:00<?, ?it/s]
loss at 0 => 463.223602: 0% | 0/150000 [00:48<?, ?it/s]
loss at 50000 => 0.610804: 33% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^N      | 5000
1/150000 [1:37:45<2:51:23, 9.72it/s]
loss at 100000 => 0.401484: 67% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^HâM-^VM-^K | 100001/150000 [3:11:07<1:23:43, 9.95it/s]
loss at 149999 => 0.378595: 100% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^HâM-^VM-^HâM-^VM-^VM-^H | 150000/150000 [4:45:29<00:00, 8.7
6it/s]
INFO:hw04.training:^[[2m2025-10-09T11:26.22.316114Z^[[0m [[[[32m^[[1m[info      ^
[0m] ^[[1mdone training      ^[[0m [[[[0m^[[1m^[[34mhw04.training^[[0
m]^[[0m
INFO:hw04:[2m2025-10-09T11:26.317243Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
mtraining concluded      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m
INFO:hw04:[2m2025-10-09T11:26.22.605988Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
msaved most recent checkpoint ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m
INFO:hw04.test:^[[2m2025-10-09T11:27.09.263316Z^[[0m [[[[32m^[[1m[info      ^
[[0m
  ^[[1mEvaluated validation set: 10000 samples, 8740 correct, Accuracy=87.40%^[[0
m [[[[0m^[[1m^[[34mhw04.test^[[0m^[[0m
INFO:hw04:[2m2025-10-09T11:27.09.263809Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
mvalidation accuracy:      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m ^[[36mval_
accuracy'^[[0m='[[35m0.874%^[[0m
INFO:hw04.test:^[[2m2025-10-09T11:27.11.177033Z^[[0m [[[[32m^[[1m[info      ^
[[0m
  ^[[1mEvaluated test set: 10000 samples, 8703 correct, Accuracy=87.03%^[[0m [[[
0m^[[1m^[[34mhw04.test^[[0m^[[0m
INFO:hw04:[2m2025-10-09T11:27.11.177436Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
mreal test accuracy:      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m^[[0m ^[[36maccu
racy'^[[0m='[[35m0.8703%^[[0m
INFO:hw04:[2m2025-10-09T11:27.11.177541Z^[[0m [[[[32m^[[1m[info      ^[[0m ^[[1
myippeeeee, close enough to state of the art i give up^[[0m [[[[0m^[[1m^[[34mhw0
4^[[0m]^[[0m

```

Oct 09, 2025 9:03

**output-cifar100-final-final.txt**

Page 1/1

```

==> Running homework 'hw04'...
/mnt/c/Users/joshu/OneDrive/Cooper Union/Deep Learning/hw04/.venv/lib/python3.13
/site-packages/pydantic/main.py:463: UserWarning: Pydantic serializer warnings:
  PydanticSerializationUnexpectedValue(Expected 'int' - serialized value may not
  be as expected [input_value=0.05, input_type=float])
  return self._pydantic_serializer_.to_python(
INFO:hw04:[2m2025-10-09T00:53:22.982046Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mSettings loaded          ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m ^[[36msett
ings'^[[0m='[[35m{'debug': False, 'random_seed': 427, 'data': {'train_ratio': 0.8
}, 'training': {'batch_size': 256, 'num_iters': 175000, 'learning_rate': 0.00015
, 'L2_weight': 0.00015}, 'model': {'input_depth': 3, 'num_classes': 100, 'layer_
depth': [64, 128, 256], 'layer_kernel_sizes': [[3, 3], [3, 3], [3, 3]], 'strides
': [1, 1, 1], 'shape': [32, 32], 'noise_std': 0.05}}^[[0m
INFO:hw04:[2m2025-10-09T00:53:28.617587Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mcifar100 data retrieved      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m ^[[36mshap
e^[[0m='[[35m(40000, 32, 32, 3)^[[0m
INFO:hw04.training:^[[2m2025-10-09T00:53:48.260701Z^[[0m [[[[32m^[[1minfo      ^[
[0m] ^[[1mstarted training      ^[[0m [[[[0m^[[1m^[[34mhw04.training^[[0
m]^[[0m

 0% | 0/175000 [00:00<?, ?it/s]
loss at 0 => 465.868805: 0% | 0/175000 [00:44<?, ?it/s]
loss at 50000 => 1.318601: 29% | âM-^VM-^HâM-^VM-^HâM-^VM-^J | 50001/175000
[1:36:43<3:54:45, 8.87it/s]
loss at 100000 => 0.769379: 57% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^K | 100001/175000 [3:12:34<2:10:13, 9.60it/s]
loss at 150000 => 0.637122: 86% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^HâM-^VM-^HâM-^VM-^L | 150001/175000 [4:47:41<41:22, 10.07it/s]
loss at 174964 => 0.592589: 100% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^I | 174965/175000 [5:35:21<00:03, 9.7
6it/s]
loss at 174999 => 0.812800: 100% | âM-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^Hâ
M-^VM-^HâM-^VM-^HâM-^VM-^HâM-^VM-^H | 175000/175000 [5:35:28<00:00, 8.6
9it/s]
INFO:hw04.training:^[[2m2025-10-09T06:29:16.912630Z^[[0m [[[[32m^[[1minfo      ^[
[0m] ^[[1mdone training      ^[[0m [[[[0m^[[1m^[[34mhw04.training^[[0
m]^[[0m
INFO:hw04:^[[2m2025-10-09T06:29:16.921884Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mtraining concluded      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m
INFO:hw04:^[[2m2025-10-09T06:29:17.291625Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
msaved most recent checkpoint      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m
INFO:hw04.test:^[[2m2025-10-09T06:30:04.425090Z^[[0m [[[[32m^[[1minfo      ^[[0m]
^[[1mEvaluated validation set: 10000 samples, 8167 correct (top-5), Accuracy=81
.67%^[[0m [[[[0m^[[1m^[[34mhw04.test^[[0m]^[[0m
INFO:hw04:^[[2m2025-10-09T06:30:04.425528Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mvalidation top5 accuracy:      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m ^[[36mval_
accuracy^[[0m='[[35m0.8167^[[0m
INFO:hw04.test:^[[2m2025-10-09T06:30:06.457077Z^[[0m [[[[32m^[[1minfo      ^[[0m]
^[[1mEvaluated test set: 10000 samples, 8134 correct (top-5), Accuracy=81.34%^[[0
m [[[[0m^[[1m^[[34mhw04.test^[[0m]^[[0m
INFO:hw04:^[[2m2025-10-09T06:30:06.457469Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mreal test top5 accuracy:      ^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m ^[[36maccu
racy^[[0m='[[35m0.8134^[[0m
INFO:hw04:^[[2m2025-10-09T06:30:06.457569Z^[[0m [[[[32m^[[1minfo      ^[[0m] ^[[1
mthis is not state of the art :(^[[0m [[[[0m^[[1m^[[34mhw04^[[0m]^[[0m

```

Oct 09, 2025 9:33

**project\_summary.txt**

Page 1/1

I started by literally training on cifar10 without changing anything meaning i didnt use data augmentation, add group normalization, or residual blocks. this was purely a CNN w 2 convolutional layers followed by fully connected layers and i got like 65% which isn't bad for relying almost solely on the inductive biases of a convolution. It trained quickly but obviously didn't generalize.

I then added data augmentation with crop, color jitter, flips(vertical and horizontal), and 90 degree rotations which yielded a cifar10 accuracy of roughly 80%. Adding group normalization made it less intensive for my GPU to train larger batch sizes faster.

Additionally, i used the residual blocks from the nnx module to implement skip connections. i used kaiming he initialization bc the PReLU paper said it was good lol.

All of this got me to an accuracy of like 92% on the validation set which was my best score and better than what i currently have lol. I then added awgn to activations during training because i thought that it would help regularize and be able to simulate data variability but

it really made no difference. I actually thought i was so smart for adding awgn in the module but i guess im just not that guy. After hitting a plateau, i tried adjusting my data augmentations by getting rid of rotation bc it was just too annoying to consider how the orientation of certain objects should be included also jax has no rotation function so unless its by 90 degrees, idk how to rotate easily. I then added awgn to the pictures to try and make the models more robust as well as greyscale, as suggested by Vaibhav. He also said to try inverting the colors and Professor, you are so right because he said that you said it makes no fucking sense bc why would it? His rational was that a green frog inverted to a black frog is still a frog???? first of all inverting green is not black, its magenta. and second of all, most animals come only in certain colors??? Anyway, sorry for the side tangent. It must have been Vaibhav's bad luck because my model declined in performance after including awgn and greyscale to the data augmentation to 88% on the validation set so that's where I'm at now. It's not state of the art. Every paper boasts 99% on CIFAR10 but i found some random guy boasting his 93% which i will interpret as state of the art bc he is all but one man and not the DL avengers assembled to conquer CIFAR10, CIFAR100, and every possible dataset.

It's worth noting that for the trial runs, i trained until the loss started to converge to a number around 1 which took about 15k iterations and 20 minutes. I really thought increasing that number to like 200k would help because the loss would drop to like .2 in some cases but i only got a small increase in performance of a couple percent meaning that i probably reached a point where i was overfitting.

Sorry that i couldn't even get 90 percent on CIFAR100 top5 accuracy btw despite even my extension. I tried; I really did. Thanks for the extension though.

Oct 08, 2025 0:35

test.py

Page 1/2

```

import jax.numpy as jnp
import numpy as np
import structlog
import optax

from .model import Classify
from .data import cifar10

log = structlog.get_logger()

def test_accuracy(
    model: Classify,
    data: cifar10,
    batch_size: int,
    validation: bool = True
) -> float:
    if validation == True:
        x_np, y_np = data.get_validation()
    else:
        x_np, y_np = data.get_test()

    num_samples = x_np.shape[0]

    # ensure labels are 1-D integer arrays
    y_np = np.asarray(y_np).reshape(-1)

    #batch predictions
    preds = []
    for start in range(0, num_samples, batch_size):
        end = start + batch_size
        xb = jnp.asarray(x_np[start:end])
        logits = model(xb, train=False)
        batch_pred = jnp.argmax(logits, axis=-1)
        preds.append(batch_pred)

    #compute accuracy
    all_preds = jnp.concatenate(preds, axis=0)
    all_labels = jnp.asarray(y_np)
    correct = jnp.sum(all_preds == all_labels)
    accuracy = correct / num_samples

    log.info(
        "Evaluated %s set: %d samples, %d correct, Accuracy=%.2f%%",
        "validation" if validation else "test",
        int(num_samples),
        int(correct),
        float(accuracy * 100),
    )

    return float(accuracy)

def test_top5(
    model: Classify,
    data: cifar10,
    batch_size: int,
    validation: bool = True
) -> float:
    """Compute top-5 accuracy on the validation or test set."""

    if validation:
        x_np, y_np = data.get_validation()
    else:
        x_np, y_np = data.get_test()

    num_samples = x_np.shape[0]

    # ensure labels are 1-D integer arrays
    y_np = np.asarray(y_np).reshape(-1)

```

Oct 08, 2025 0:35

test.py

Page 2/2

```

correct_top5 = 0

for start in range(0, num_samples, batch_size):
    end = start + batch_size
    xb = jnp.asarray(x_np[start:end])
    yb = jnp.asarray(y_np[start:end])

    logits = model(xb, train=False)

    # get indices of top 5 predictions per sample
    top5 = jnp.argsort(logits, axis=-1)[:, -5:] # shape (batch_size, 5)

    # check if true label is in top 5 predictions
    correct_batch = jnp.any(top5 == yb[:, None], axis=-1)
    correct_top5 += jnp.sum(correct_batch)

top5_accuracy = correct_top5 / num_samples

log.info(
    "Evaluated %s set: %d samples, %d correct (top-5), Accuracy=%.2f%%",
    "validation" if validation else "test",
    int(num_samples),
    int(correct_top5),
    float(top5_accuracy * 100),
)

return float(top5_accuracy)

```

Oct 08, 2025 22:37

## training.py

Page 1/4

```

import jax.numpy as jnp
import jax
import numpy as np
import optax
import structlog
import orbax.checkpoint as ocp
from typing import Optional
from pathlib import Path

from flax import nnx
from tqdm import trange
from .config import TrainingSettings
from .data import cifar10
from .model import Classify

log = structlog.get_logger()

@nnx.jit
def train_step_cifar10(
    model: Classify,
    optimizer: nnx.Optimizer,
    x: jnp.ndarray,
    y: jnp.ndarray,
):
    """Perform on training step"""

    # x shape: (batch_size, 32, 32, 3); y shape: (batch_size,)
    def loss_function(model: Classify):
        logits = model(x, True)
        ce_loss = jnp.mean(optax.softmax_cross_entropy_with_integer_labels(logits, y))
        L2_loss = model.L2_loss()
        return ce_loss + L2_loss

    loss, grads = nnx.value_and_grad(loss_function)(model)
    optimizer.update(model, grads)
    return loss

def augment(
    key: jnp.array = jax.random.PRNGKey(472),
    images: jnp.array = jnp.ndarray,
    *,
    p_hflip: float = 0.35, #probabilities for performing these augmentations
    pad: int = 4,
    brightness: float = 0.2, #changes in brightness, contrast, and saturation
    contrast: float = 0.2,
    saturation: float = 0.2,
    p_cutout: float = 0.35,
    # i define this in the cutout function bc i break vmap otherwise and im too
    # lazy to figure out why: max_cutout: int = 8,
    p_greyscale: float = 0.3,
    p_noise: float = 0.85,
    noise_std: float = 0.15, # stddev of additive Gaussian noise
) -> jnp.ndarray:
    """
    Batch augmentation
    images: float32 in [0,1], shape (B,H,W,C).
    Returns augmented images, same shape and dtype.
    """
    B, H, W, C = images.shape

    # Split top-level key into independent subkeys.
    # We need separate keys for flips, crop, color, cutout, greyscale, noise мас
    k, noise_rng =
        key, k_flip_h, k_crop, k_color, k_cutout, k_grey, k_noise_mask, k_noise =
            (jax.random.split(key, 8)
    )

```

Oct 08, 2025 22:37

## training.py

Page 2/4

```

    """i got rid of rotation because many classes in the cifar datasets dont make sense to flip and rotating
    it by a small amount instead didnt really help much. also it probably didnt help that i couldnt find a rotation funciton i
    n the
    jax documentation lol"""

    # Horizontal flip
    hmask = jax.random.bernoulli(k_flip_h, p_hflip, (B,))
    images = jnp.where(hmask[:, None, None], images[:, :, ::-1, :], images)

    """i also axed vertical flip because its genuinely just not worth it given that many of the classes in both cifar data
    sets
    #wouldnt realistically have many vertically flipped images"""

    # Crop (random crop from reflected padding)
    padded = jnp.pad(images, ((0, 0), (pad, pad), (pad, pad), (0, 0)), mode="reflect")
    max_x = padded.shape[1] - H
    max_y = padded.shape[2] - W
    k_crop_x, k_crop_y = jax.random.split(k_crop, 2)
    xs = jax.random.randint(k_crop_x, (B,), 0, max_x + 1)
    ys = jax.random.randint(k_crop_y, (B,), 0, max_y + 1)

    def crop_one(img, x, y):
        return jax.lax.dynamic_slice(img, (x, y, 0), (H, W, C))

    images = jax.vmap(crop_one)(padded, xs, ys)

    # Color jitter (brightness / contrast / saturation)
    k_b, k_c, k_s = jax.random.split(k_color, 3)
    bright_f = jax.random.uniform(k_b, (B,), minval=-brightness, maxval=brightness)
    contrast_f = jax.random.uniform(k_c, (B,), minval=1.0 - contrast, maxval=1.0 + contrast)
    sat_f = jax.random.uniform(k_s, (B,), minval=-saturation, maxval=saturation)

    def apply_color(bf, cf, sf, img):
        # brightness
        img = img + bf
        # contrast
        mean = jnp.mean(img, axis=(0, 1), keepdims=True)
        img = (img - mean) * cf[..., None, None] + mean
        # saturation (convert to luminance and lerp)
        lum = img[..., 0] * 0.2989 + img[..., 1] * 0.5870 + img[..., 2] * 0.1140
        lum = lum[..., None]
        img = img * (1.0 - sf)[..., None, None] + lum * sf[..., None, None]
        return jnp.clip(img, 0.0, 1.0)

    images = jax.vmap(apply_color)(bright_f, contrast_f, sat_f, images)

    # Cutout
    cut_mask = jax.random.bernoulli(k_cutout, p_cutout, (B,)) # which images get cutout
    # create 4 keys per image (one for cut_h, cut_w, ky, kx)
    per_cut_keys = jax.random.split(k_cutout, B * 4).reshape((B, 4, 2))

    def do_cutout(karr, do, img, max_cutout = 8):
        Hc, Wc, Cc = img.shape

        def true_fn(args):
            karr, img = args
            k0, k1, k2, k3 = karr

            # pick random cutout size
            cut_h = jax.random.randint(k0, (0, 1, max_cutout + 1)
            cut_w = jax.random.randint(k1, (0, 1, max_cutout + 1)

            # pick random top-left corner

```

Oct 08, 2025 22:37

**training.py**

Page 3/4

```

ky = jax.random.randint(k2, (), 0, Hc - cut_h + 1)
kx = jax.random.randint(k3, (), 0, Wc - cut_w + 1)

# create a mask of ones (static shape)
mask = jnp.ones((Hc, Wc, Cc), dtype=img.dtype)

# create indices for cutout
rows = jnp.arange(Hc)
cols = jnp.arange(Wc)
row_mask = (rows < ky) | (rows >= ky + cut_h)
col_mask = (cols < kx) | (cols >= kx + cut_w)

# broadcast row/col masks
cut_mask_local = jnp.outer(row_mask, col_mask).astype(img.dtype)
cut_mask_local = cut_mask_local[..., None] # expand to (H, W, 1)
mask = mask * cut_mask_local # apply mask

return img * mask

def false_fn(args):
    return args[1]

return jax.lax.cond(do, true_fn, false_fn, (karr, img))

images = jax.vmap(do_cutout)(per_cut_keys, cut_mask, images)

#greyscale
if p_greyscale > 0.0:
    grey_mask = jax.random.bernoulli(k_grey, p_greyscale, (B,)) # which images -> greyscale
    # compute luminance and stack to 3 channels for those images
    lum = images[..., 0] * 0.2989 + images[..., 1] * 0.5870 + images[..., 2] * 0.1140 # (B, H, W)
    lum = lum[..., None] # (B, H, W, 1)
    grey_images = jnp.concatenate([lum, lum, lum], axis=-1) # (B, H, W, 3)
    images = jnp.where(grey_mask[:, None, None, None], grey_images, images)

#awgn
if p_noise > 0.0 and noise_std > 0.0:
    noise_mask = jax.random.bernoulli(k_noise_mask, p_noise, (B,))
    noise = jax.random.normal(k_noise, (B, H, W, C)) * noise_std
    noisy = images + noise
    noisy = jnp.clip(noisy, 0.0, 1.0)
    images = jnp.where(noise_mask[:, None, None, None], noisy, images)

return jnp.clip(images, 0.0, 1.0)

def cifar10_train(
    model: Classify,
    optimizer: nnx.Optimizer,
    data: cifar10,
    settings: TrainingSettings,
    np_rng: np.random.Generator,
) -> None:
    """train with grad descend"""
    log.info("started training")

    # create a persistent jnp PRNGKey seeded from the numpy RNG you already have
    jnp_key = jax.random.PRNGKey(int(np_rng.integers(0, 2 ** 31 - 1)))

    bar = trange(settings.num_iters)
    for i in bar:
        x_np, y_np = data.get_batch(np_rng, settings.batch_size)
        x = jnp.asarray(x_np, dtype=jnp.float32)
        y = jnp.asarray(y_np, dtype=jnp.int32).reshape(-1)

        # get subkey and augment the batch (augmentation is outside the jitted train step)
        jnp_key, subkey = jax.random.split(jnp_key)

```

Oct 08, 2025 22:37

**training.py**

Page 4/4

```

x_aug = augment(subkey, x)

loss = train_step_cifar10(model, optimizer, x_aug, y)

bar.set_description(f"loss at {i} => {loss:.6f}")
bar.refresh()

log.info("done training")

```

Oct 09, 2025 2:35

\_\_init\_\_.py

Page 1/3

```

import jax
import numpy as np
import structlog
import optax
from pathlib import Path
import orbax.checkpoint as ocp
from flax import nnx

from .config import load_settings
from .data import cifar10
from .model import Classify
from .logging import configure_logging
from .training import cifar10_train
from .test import test_accuracy, test_top5

def main() -> None:
    """CLI entry point."""
    settings = load_settings()
    configure_logging()
    log = structlog.get_logger()
    log.info("Settings loaded", settings=settings.model_dump())

    # JAX PRNG
    key = jax.random.PRNGKey(472) #472 for ECE 472
    data_key, model_key = jax.random.split(key)
    np_rng = np.random.default_rng(np.array(data_key))

    data = cifar10(
        rng = np_rng,
        train_ratio = settings.data.train_ratio,
    )
    if settings.model.num_classes == 10:
        log.info("cifar10 data retrieved", shape = data.x_train.shape)
    else:
        log.info("cifar100 data retrieved", shape = data.x_train.shape)

    model = Classify(
        rngs = nnx.Rngs(params=model_key),
        num_classes = settings.model.num_classes,
        shape = settings.model.shape,
        input_depth = settings.model.input_depth,
        layer_depths = settings.model.layer_depth,
        layer_kernel_sizes = settings.model.layer_kernel_sizes,
        strides = settings.model.strides,
        L2_weight = settings.training.L2_weight,
    )

    optimizer_schedule = optax.cosine_decay_schedule(settings.training.learning_rate, settings.training.num_iters,) #same schedule as you
    optimizer = nnx.Optimizer(
        model, optax.adam(optimizer_schedule), wrt=nnx.Param
    )

    cifar10_train(model, optimizer, data, settings.training, np_rng)
    log.info("training concluded")

    ckpt_dir = ocp.test_utils.erase_and_create_empty("/saves")
    graphdef, state = nnx.split(model)
    checkpoint = ocp.StandardCheckpointer()
    checkpoint.save(ckpt_dir / "cifar10", state)
    log.info("saved most recent checkpoint")

    #validation and test set and stuff
    if settings.model.num_classes == 10:
        validation_accuracy = test_accuracy(
            data = data,
            batch_size = settings.training.batch_size,
            model = model,
            validation = True,
        )

```

Oct 09, 2025 2:35

\_\_init\_\_.py

Page 2/3

```

        )
        log.info("validation accuracy:", val_accuracy = validation_accuracy)
        real_test_accuracy = test_accuracy(
            data = data,
            batch_size = settings.training.batch_size,
            model = model,
            validation = False,
        )
        log.info("real test accuracy:", accuracy = real_test_accuracy)

    else:
        validation_accuracy = test_top5(
            data = data,
            batch_size = settings.training.batch_size,
            model = model,
            validation = True,
        )
        log.info("validation top5 accuracy:", val_accuracy = validation_accuracy)
        real_test_accuracy = test_top5(
            data = data,
            batch_size = settings.training.batch_size,
            model = model,
            validation = False,
        )
        log.info("real test top5 accuracy:", accuracy = real_test_accuracy)

    if real_test_accuracy >= 0.85:
        log.info("yipppeeee, close enough to state of the art i give up")
    else:
        log.info("this is not state of the art :(")

def test10():
    settings = load_settings()
    configure_logging()
    log = structlog.get_logger()
    log.info("Settings loaded", settings=settings.model_dump())

    # JAX PRNG
    key = jax.random.PRNGKey(472) #472 for ECE 472
    data_key, model_key = jax.random.split(key)
    np_rng = np.random.default_rng(np.array(data_key))

    log.info("evaluating on test set for CIFAR10!!")

    model = Classify(
        rngs = nnx.Rngs(params=model_key),
        num_classes = settings.model.num_classes,
        shape = settings.model.shape,
        input_depth = settings.model.input_depth,
        layer_depths = settings.model.layer_depth,
        layer_kernel_sizes = settings.model.layer_kernel_sizes,
        strides = settings.model.strides,
        L2_weight = settings.training.L2_weight,
    )

    data = cifar10(
        rng = np_rng,
        train_ratio = settings.data.train_ratio,
    )

    ckpt_dir = Path("/saves").resolve()
    checkpoint = ocp.StandardCheckpointer()
    graphdef, state = nnx.split(model)
    state = jax.tree_util.tree_map(ocp.utils.to_shape_dtype_struct, state)
    restored = checkpoint.restore(ckpt_dir / 'most_recent', state)
    model = nnx.merge(graphdef, restored)
    log.info("restored model from checkpoint", model = model)

    real_test_accuracy = test_accuracy(

```

Oct 09, 2025 2:35

\_\_init\_\_.py

Page 3/3

```
data = data,
batch_size = settings.training.batch_size,
model = model,
validation = False,
)
log.info("real test accuracy:", accuracy = real_test_accuracy)

def test100():
    settings = load_settings()
    configure_logging()
    log = structlog.get_logger()
    log.info("Settings loaded", settings=settings.model_dump())

    # JAX PRNG
    key = jax.random.PRNGKey(472) #472 for ECE 472
    data_key, model_key = jax.random.split(key)
    np_rng = np.random.default_rng(np.array(data_key))

    log.info("evaluating on test set for CIFAR100!!")

    model = Classify(
        rngs = nnx.Rngs(params=model_key),
        num_classes = settings.model.num_classes,
        shape = settings.model.shape,
        input_depth = settings.model.input_depth,
        layer_depths = settings.model.layer_depth,
        layer_kernel_sizes = settings.model.layer_kernel_sizes,
        strides = settings.model.strides,
        L2_weight = settings.training.L2_weight,
    )
    data = cifar10( #should be cifar100 but im too lazy to change
        rng = np_rng,
        train_ratio = settings.data.train_ratio,
    )
    ckpt_dir = Path("/saves").resolve()
    checkpointer = ocp.StandardCheckpointer()
    graphdef, state = nnx.split(model)
    state = jax.tree_util.tree_map(ocp.utils.to_shape_dtype_struct, state)
    restored = checkpointer.restore(ckpt_dir/ 'most_recent', state)
    model = nnx.merge(graphdef, restored)
    log.info("restored model from checkpoint", model = model)

    real_test_accuracy = test_accuracy(
        data = data,
        batch_size = settings.training.batch_size,
        model = model,
        validation = False,
    )
    log.info("real test accuracy:", accuracy = real_test_accuracy)
```